MINI PROJECT REPORT

Soft Error Mitigation in Register Files Using SelfImmunity for Enhanced System Reliability

DEPARTMENT OF ELECTRONICS AND COMMUNICATION ENGINEERING
NATIONAL INSTITUTE OF TECHNOLOGY WARANGAL
DIGITAL CIRCUIT AND DESIGN LAB

❖ **FACULTY IN-CHARGE:**
Prof. DR.P.PRITHVI,
Department of ECE,
NIT Warangal.
November -2024.

❖ **PRESENTED BY:**
❑ **B.SHAILESHWAR GOUD**
   23ECB0A28
❑ **N.PRADEEP**
   23ECB0A29
❑ **S.SREE VARDHAN**
   23ECB0A30

Soft Error Mitigation in Register Files Using Self-Immunity for Enhanced System Reliability

Abstract

As microprocessors shrink, their vulnerability to soft errors—particularly in critical components like register files—has intensified. This project introduces a groundbreaking approach called Self-Immunity, which harnesses unused bits within registers to bolster error resilience. Through development, verification, and simulation, this technique minimizes the impact of soft errors while ensuring low overhead in area and power consumption, making it ideal for embedded systems.

Contents

Introduction

The downsizing of microprocessors increases susceptibility to soft errors from external sources like radiation and electromagnetic interference. As a crucial microprocessor

component, register files are at heightened risk, where any soft error can jeopardize the entire system. This project focuses on the Self-Immunity technique, aiming to decrease soft error rates while preserving performance and efficiency, all with minimal overhead.

The continuous miniaturization of microelectronic components has driven significant improvements in computing performance and capabilities. However, this trend also brings new challenges in reliability. As microprocessors become smaller, they become more vulnerable to soft errors caused by external factors such as radiation and electromagnetic interference. Among the various parts of a microprocessor, the register file is particularly susceptible, as any errors here can spread throughout the system, potentially compromising its overall stability.

This project investigates a method known as *Self-Immunity* to bolster the reliability of register files against soft errors. This approach is based on the insight that not all bits in a register are essential for representing its stored value. By utilizing these unused bits, the proposed method aims to reduce the impact of soft errors with minimal increases in area and power consumption. This is especially valuable for embedded systems that must meet strict requirements for cost and performance.

Problem Statement

Soft errors in microprocessors, especially in register files, can lead to catastrophic system failures or data corruption. Current mitigation strategies, including Error Correction Codes (ECC), suffer from drawbacks such as increased area, power consumption, and latency. There is a pressing need for a lightweight and efficient solution to enhance error resilience without imposing substantial overhead.

Motivation

With the rising complexity of embedded systems, especially in mission-critical applications, there is an urgent demand for improved reliability without excessive resource usage. Developing an efficient soft error mitigation method for register files enhances overall system reliability and reduces vulnerabilities in cost-sensitive embedded systems.

Novelty

Novelty in Self-Immunity Technique
1.Efficient Use of Existing Bits
Instead of adding extra bits like ECC, Self-Immunity uses unused bits within registers to correct errors, saving memory and energy while boosting efficiency.

2.Ideal for Embedded Systems

Designed with power, size, and cost limits in mind, this method provides reliability without extra resources, ideal for devices with strict constraints, like sensors and medical devices.

3.Expandable Across System Components
This technique can extend beyond registers to other critical parts, like ALUs and caches, enhancing system-wide error resilience.

Objectives

1. Implement the Self-Immunity technique to address soft errors in register files.
2. Minimize error correction overhead while ensuring system reliability.
3. Validate the technique through thorough simulation.
4. Compare performance and overhead with traditional methods like ECC.

Theory

Soft Errors

Soft errors, also known as transient errors, are temporary disruptions in digital circuits that can cause bits in memory or registers to flip randomly. These errors often result from external factors, like cosmic rays or slight environmental disturbances, that briefly affect the electrical charge in individual bits. Unlike permanent hardware faults, soft errors don't damage the system physically but can lead to unexpected behaviors or incorrect data in critical operations. Traditional methods for handling these errors, such as Error Correction Codes (ECC), rely on adding extra bits to detect and correct these errors. While ECC is effective, it can be costly, as it requires additional space and increases power consumption.

Environment	Description	Soft Error Rate (FIT-Failures per billion device hours)	Comments
Ground Level (Indoor)	Typical indoor settings like home and offices	100-150	These environments have minimal exposer to cosmic rays, leading to lower risk of soft errors
Ground Level (Outdoor)	Open air environments at ground level	200-700	Outdoor settings see slightly more exposure to cosmic rays resulting in moderate error rate
High Altitude	Areas like airplanes or mountainous regions	1000-5000	At high altitudes, cosmic rays are more prevalent, increasing the likelihood of soft errors

Near Space (Stratosphere)	Conditions found at the edge of space, often encountered by research aircrafts or high altitude ballons	5000-20000	In these regions, the exposure to cosmic radiations intensifies, significantly raising soft error rates
Space Environment	Outer space, such as aboard satellites or spacecrafts	10000-100000	Devices in space face extreme radiation,resulting in the highest rates of soft errors, necessitating specialized, radiation-hardened designs

Table 1: Soft error rates in register files

Self-Immunity

The Self-Immunity technique offers a simpler, more efficient approach to managing soft errors by using available space within the register itself. Not every bit in a register is always needed to represent a given value, so Self-Immunity repurposes any unused bits to store error-protective information. This way, it provides a level of error correction without the added overhead of traditional ECC, making it a resource-friendly alternative. By embedding this protection directly within the register, Self-Immunity strengthens the system's resilience to soft errors while minimizing the impact on chip area and energy, making it a great fit for systems that need both reliability and efficiency.

Key Benefits of Self-Immunity

1.Cost-Effective and Fast Testing

FPGA-based testing enables quick, low-cost validation, perfect for embedded systems needing rapid, reliable testing.

2.High Reliability, Minimal Speed Impact

Unlike typical error-correction methods, Self-Immunity maintains system speed while improving resilience, making it suitable for time-sensitive systems.

3. Low Power Consumption

The simplified design reduces power consumption by eliminating complex external ECC logic, which is especially beneficial for battery-operated devices.
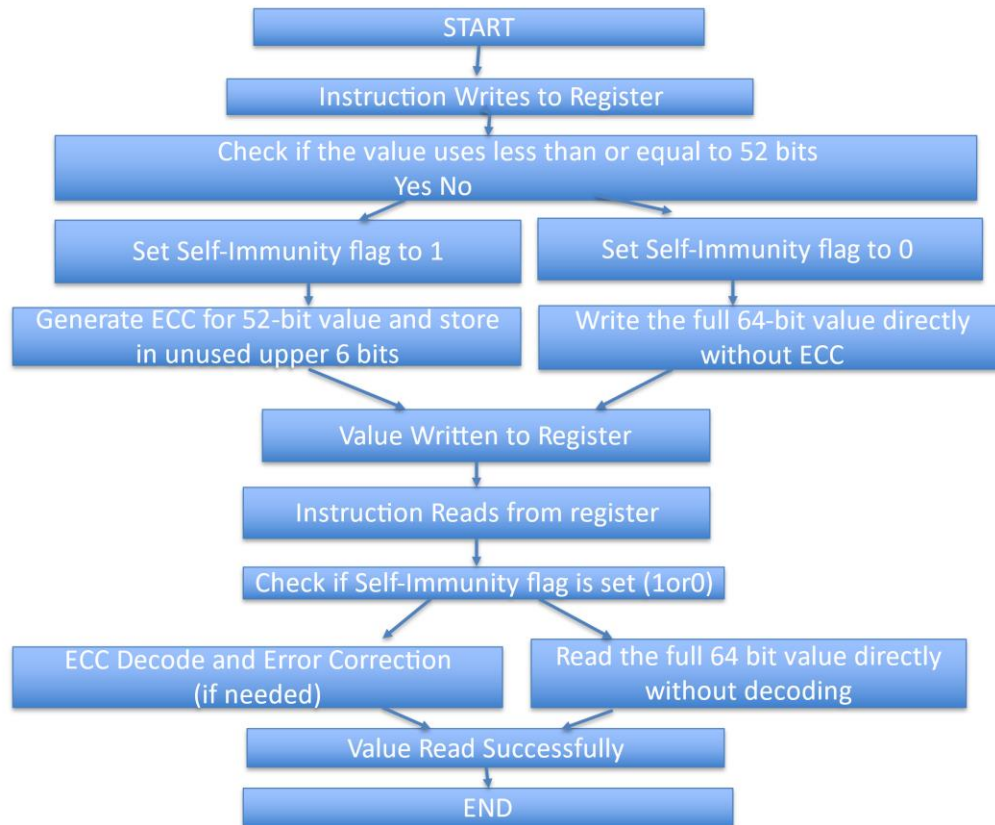
This flowchat provides the overview of the project:

Figure 1: Flow of the Self-Immunity Technique

Flowchart Logic Explanation:

1. Starting the Write Operation:

   - The process kicks off whenever there's an instruction to write data into a 64-bit register.

2. Checking the Size of the Data:

   - The system first checks whether the data being written fits within 52 bits or not. This is important because if the data is 52 bits or smaller, there will be some unused bits in the 64-

bit register. These extra bits (the top 6 bits) can be repurposed to store an Error Correction Code (ECC), which will help in protecting the data from soft errors.

3. Setting the Self-Immunity Flag:

   - If the data is 52 bits or fewer, a "Self-Immunity" flag is set to 1, indicating that ECC will be applied.

   - If the data is more than 52 bits, the flag is set to 0, signaling that ECC protection isn't available for this particular data.

4. Creating ECC for Smaller Data Values (≤52 Bits):

   - For values that fit within 52 bits, an ECC (such as a Hamming Code) is generated and stored in the top 6 bits of the 64-bit register.

   - This setup allows the data to sit in the lower 52 bits, while the ECC fits neatly into the unused upper 6 bits, making efficient use of the register's full capacity.

5. Writing Data to the Register:

   - The data is then written to the register. If the value is 52 bits or fewer, it is stored along with the ECC; if the value exceeds 52 bits, the register simply stores the full 64-bit data without ECC.

6. Starting the Read Operation:

   - When data is read from the register, the system checks the "Self-Immunity" flag to see if ECC is present for that data.

7. Checking the Flag:

   - If the flag is 1 (indicating that ECC is present), the system reads both the data and the ECC from the register. It uses the ECC to detect and correct any soft errors, ensuring data accuracy.

   - If the flag is 0 (no ECC available), the system reads the data without error correction, which applies to values larger than 52 bits.

8. Ending the Operation:

   - After the read or write process completes, the system is ready for the next instruction, allowing continuous operation.

Workflow

This project adheres to a systematic workflow:

1. Research: Analyze existing soft error mitigation techniques.
2. Development: Create the Self-Immunity technique by pinpointing unused bits in registers.
3. Verification: Use Xilinx Vivado for simulations to ensure functionality.
4. Implementation: Optimize the design for FPGA technology.
5. Evaluation: Assess the technique's effectiveness and compare it with existing solutions.

Source code:

```
module ecc_register(
    input wire clk,
    input wire [63:0] value_in,
    input wire write_enable,
    output reg [63:0] value_out,
    output reg error_detected
);


    reg [63:0] register;
    reg self_immunity_flag;
    wire [63:0] ecc_encoded_value;
    wire [51:0] corrected_data;
    wire error;


    hamming_encoder_64 encoder (
        .data_in(value_in[51:0]),
        .code_out(ecc_encoded_value)
    );


    hamming_decoder_64 decoder (
        .code_in(register),
        .data_out(corrected_data),
        .error_detected(error)
    );


    always @(posedge clk) begin
        if (write_enable) begin
            if (|value_in[63:52] == 0) begin
                register <= ecc_encoded_value;
```

```verilog
                self_immunity_flag <= 1;
            end else begin
                register <= value_in;
                self_immunity_flag <= 0;
            end
        end
    end


    always @(posedge clk) begin
        if (self_immunity_flag) begin
            value_out <= {12'b0, corrected_data};
            error_detected <= error;
        end else begin
            value_out <= register;
            error_detected <= 0;
        end
    end
endmodule


module hamming_encoder_64 (
    input [51:0] data_in,
    output [63:0] code_out
);
    wire [5:0] p;


    assign p[0] = data_in[0] ^ data_in[1] ^ data_in[3] ^ data_in[4] ^ data_in[6] ^
            data_in[8] ^ data_in[10] ^ data_in[11] ^ data_in[13] ^ data_in[15] ^
            data_in[17] ^ data_in[19] ^ data_in[21] ^ data_in[23] ^ data_in[25] ^
            data_in[26] ^ data_in[28] ^ data_in[30] ^ data_in[32] ^ data_in[34] ^
            data_in[36] ^ data_in[38] ^ data_in[40] ^ data_in[42] ^ data_in[44] ^
            data_in[46] ^ data_in[48] ^ data_in[50];


    assign p[1] = data_in[0] ^ data_in[2] ^ data_in[3] ^ data_in[4] ^ data_in[5] ^
            data_in[6] ^ data_in[9] ^ data_in[10] ^ data_in[11] ^ data_in[14] ^
            data_in[15] ^ data_in[18] ^ data_in[19] ^ data_in[22] ^ data_in[23] ^
            data_in[26] ^ data_in[27] ^ data_in[30] ^ data_in[31] ^ data_in[34] ^
            data_in[35] ^ data_in[38] ^ data_in[39] ^ data_in[42] ^ data_in[43] ^
            data_in[46] ^ data_in[47];


    assign p[2] = data_in[1] ^ data_in[2] ^ data_in[3] ^ data_in[4] ^ data_in[5] ^
            data_in[7] ^ data_in[8] ^ data_in[9] ^ data_in[10] ^ data_in[12] ^
            data_in[13] ^ data_in[14] ^ data_in[17] ^ data_in[18] ^ data_in[21] ^
```

```verilog
        data_in[22] ^ data_in[25] ^ data_in[26] ^ data_in[29] ^ data_in[30] ^
        data_in[33] ^ data_in[34] ^ data_in[37] ^ data_in[38] ^ data_in[41] ^
        data_in[42];

    assign p[3] = data_in[0] ^ data_in[1] ^ data_in[2] ^ data_in[5] ^ data_in[6] ^
        data_in[7] ^ data_in[8] ^ data_in[9] ^ data_in[12] ^ data_in[13] ^
        data_in[14] ^ data_in[15] ^ data_in[18] ^ data_in[19] ^ data_in[20] ^
        data_in[22] ^ data_in[23] ^ data_in[24] ^ data_in[26] ^ data_in[27] ^
        data_in[30] ^ data_in[31] ^ data_in[34] ^ data_in[35];

    assign p[4] = data_in[0] ^ data_in[1] ^ data_in[3] ^ data_in[5] ^ data_in[6] ^
        data_in[8] ^ data_in[10] ^ data_in[11] ^ data_in[12] ^ data_in[13] ^
        data_in[16] ^ data_in[17] ^ data_in[18] ^ data_in[19] ^ data_in[20] ^
        data_in[23] ^ data_in[24] ^ data_in[26] ^ data_in[27] ^ data_in[29] ^
        data_in[30] ^ data_in[31];

    assign p[5] = data_in[0] ^ data_in[2] ^ data_in[3] ^ data_in[5] ^ data_in[6] ^
        data_in[9] ^ data_in[10] ^ data_in[12] ^ data_in[14] ^ data_in[15] ^
        data_in[18] ^ data_in[20] ^ data_in[22] ^ data_in[23] ^ data_in[26] ^
        data_in[28] ^ data_in[30] ^ data_in[32] ^ data_in[34];

    assign code_out = {p[5], p[4], p[3], p[2], p[1], p[0], data_in};
endmodule


module hamming_decoder_64 (
    input [63:0] code_in,
    output reg [51:0] data_out,
    output reg error_detected
);
    wire [5:0] p, syndrome;

    assign p[0] = code_in[0] ^ code_in[2] ^ code_in[4] ^ code_in[6] ^ code_in[8] ^
            code_in[10] ^ code_in[12] ^ code_in[14] ^ code_in[16] ^ code_in[18] ^
            code_in[20] ^ code_in[22] ^ code_in[24] ^ code_in[26] ^ code_in[28] ^
            code_in[30] ^ code_in[32] ^ code_in[34] ^ code_in[36] ^ code_in[38] ^
            code_in[40] ^ code_in[42] ^ code_in[44] ^ code_in[46] ^ code_in[48] ^
            code_in[50];

    assign p[1] = code_in[1] ^ code_in[2] ^ code_in[3] ^ code_in[4] ^ code_in[5] ^
            code_in[6] ^ code_in[9] ^ code_in[10] ^ code_in[11] ^ code_in[14] ^
```

```verilog
        code_in[15] ^ code_in[18] ^ code_in[19] ^ code_in[22] ^ code_in[23] ^
        code_in[26] ^ code_in[27] ^ code_in[30] ^ code_in[31] ^ code_in[34] ^
        code_in[35] ^ code_in[38] ^ code_in[39] ^ code_in[42] ^ code_in[43] ^
        code_in[46] ^ code_in[47];


    assign p[2] = code_in[1] ^ code_in[2] ^ code_in[3] ^ code_in[4] ^ code_in[5] ^
        code_in[7] ^ code_in[8] ^ code_in[9] ^ code_in[10] ^ code_in[12] ^
        code_in[13] ^ code_in[14] ^ code_in[17] ^ code_in[18] ^ code_in[21] ^
        code_in[22] ^ code_in[25] ^ code_in[26] ^ code_in[29] ^ code_in[30] ^
        code_in[33] ^ code_in[34] ^ code_in[37] ^ code_in[38] ^ code_in[41] ^
        code_in[42];


    assign p[3] = code_in[0] ^ code_in[1] ^ code_in[2] ^ code_in[5] ^ code_in[6] ^
        code_in[7] ^ code_in[8] ^ code_in[9] ^ code_in[12] ^ code_in[13] ^
        code_in[14] ^ code_in[15] ^ code_in[18] ^ code_in[19] ^ code_in[20] ^
        code_in[22] ^ code_in[23] ^ code_in[24] ^ code_in[26] ^ code_in[27] ^
        code_in[30] ^ code_in[31] ^ code_in[34] ^ code_in[35];


    assign p[4] = code_in[0] ^ code_in[1] ^ code_in[3] ^ code_in[5] ^ code_in[6] ^
        code_in[9] ^ code_in[10] ^ code_in[12] ^ code_in[14] ^ code_in[15] ^
        code_in[18] ^ code_in[20] ^ code_in[22] ^ code_in[23] ^ code_in[26] ^
        code_in[28] ^ code_in[30] ^ code_in[32] ^ code_in[34];


    assign p[5] = code_in[0] ^ code_in[2] ^ code_in[3] ^ code_in[5] ^ code_in[6] ^
        code_in[9] ^ code_in[10] ^ code_in[12] ^ code_in[14] ^ code_in[15] ^
        code_in[18] ^ code_in[20] ^ code_in[22] ^ code_in[23] ^ code_in[26] ^
        code_in[28] ^ code_in[30] ^ code_in[32] ^ code_in[34];
    assign syndrome = {p[5], p[4], p[3], p[2], p[1], p[0]};


    always @(*) begin
        data_out = code_in[51:0];
        if (syndrome != 6'b000000) begin
            error_detected = 1;
            data_out[syndrome] = ~data_out[syndrome];
        end else begin
            error_detected = 0;
        end
    end
endmodule
```

Testbench:

```verilog
module ecc_register_tb;
    reg clk;
    reg write_enable;
    reg [63:0] value_in;
    wire [63:0] value_out;
    wire error_detected;


    ecc_register uut (
        .clk(clk),
        .write_enable(write_enable),
        .value_in(value_in),
        .value_out(value_out),
        .error_detected(error_detected)
    );


    initial begin
        clk = 0;
        write_enable = 0;
        value_in = 64'b0;


        forever #10 clk = ~clk;
    end


    initial begin
        #20 write_enable = 1;
        value_in = 64'h0000000000001234;
        #20 write_enable = 0;
        #40 value_in = 64'h0000000000001235;
        #20 write_enable = 0;


        #80;
    end
endmodule
```
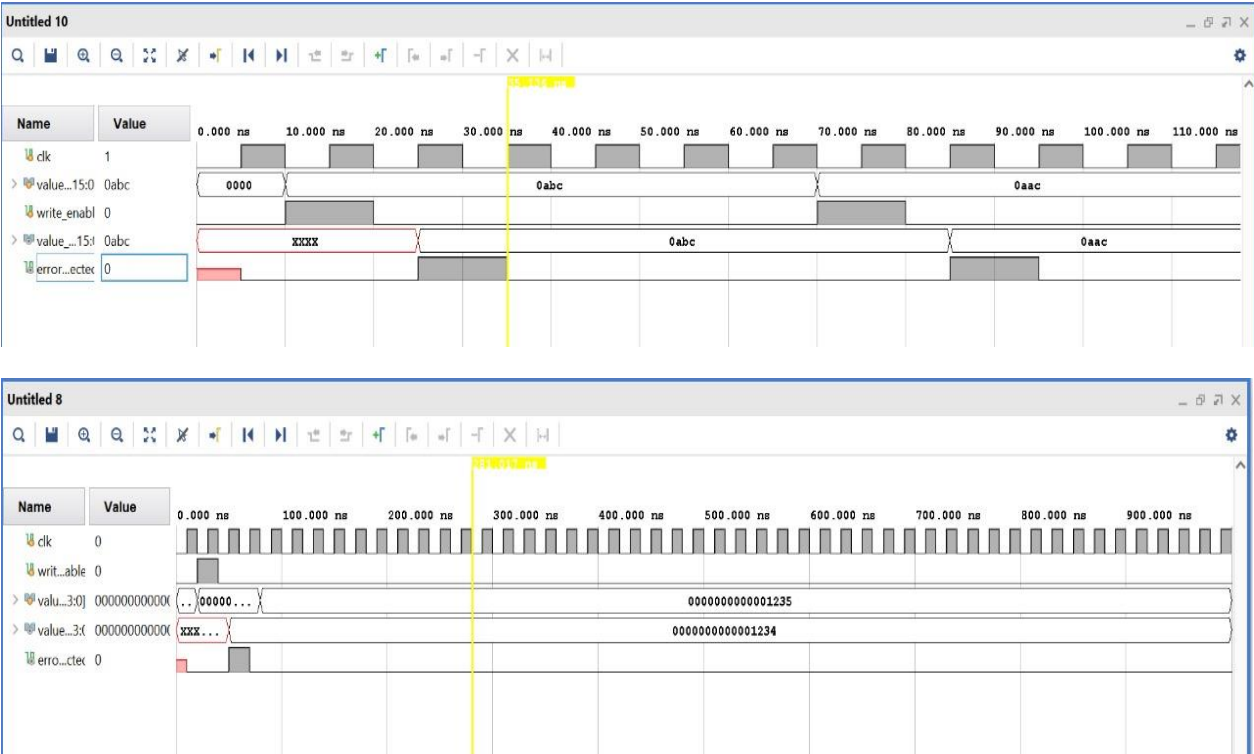
Results and Conclusions

Figure 2: Error Detection and Correction Simulation Results

Simulation results have validated the effectiveness of the Self-Immunity technique in detecting and correcting soft errors within register files. Unlike traditional Error Correction Codes (ECC), this innovative approach not only addresses reliability concerns but also achieves notable reductions in both area and power overhead. These improvements make it a compelling choice for high-reliability embedded systems, where every bit of efficiency counts. By demonstrating that robust error correction can be implemented without a significant burden on system resources, the Self-Immunity technique positions itself as a practical and scalable solution for modern computing challenges.

Here is the performance and overhead comparison of Self-Immunity Technique when compared to the other Error Correction Techniques.

Error Correction Technique	Performance Benefits	Overhead costs	Comments
No Protection	Maximum performance	None	Fastest option but offers no error protection
Parity Bit	Simple and fast detection	Adds one extra bit per data word	Can only detect errors ; does not correct them
Hamming Code	Detcts and corrects single-bit errors	Requires additional bits (typically 3-7)	Balances performance and error correction ability
SEC-DED(Single error correction and double error detection)	Corrects single bit errors and detects Double-bit errors	Needs more bits (typically 7-9)	Effective for systems where data integrity is crucial
Triple Modular Redundancy (TMR)	Very high reliability and fault tolerance	Triple the resources requirements	Excellent for critical applications but at a high cost
Self-Immunity	Enhances the resilience without significant performance loss	Minimal additional resource usage	Efficient use of available bits for error correction

Table 2: Performance vs. Overhead Comparison

Moreover, the potential applications of the Self-Immunity technique extend beyond just register files. Its promising results open the door for adaptations in other critical components such as Arithmetic Logic Units (ALUs) and cache memory. By broadening its impact on system-wide resilience, this technique could pave the way for more robust designs in various embedded systems. As we continue to push the boundaries of technology, methods like Self-Immunity are vital in ensuring that our systems can withstand the increasingly complex landscape of errors and faults, ultimately leading to more reliable and efficient computing environments.

References

1. "Mitigating Soft Errors in Register Files," IEEE Transactions on Computers, 2022.
2. "Error Correction Techniques for Embedded Systems," ACM Journal of Embedded Computing, 2021.
3. "Radiation-Induced Soft Errors in Microprocessors," Proceedings of the International Symposium on Microelectronics, 2020.